


A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

# Attacking the Heap

By Joshua Wong

# Why should you learn heap exploitation?

The root cause of exploited vulnerabilities provide hints on attacker preference & ease of exploitability

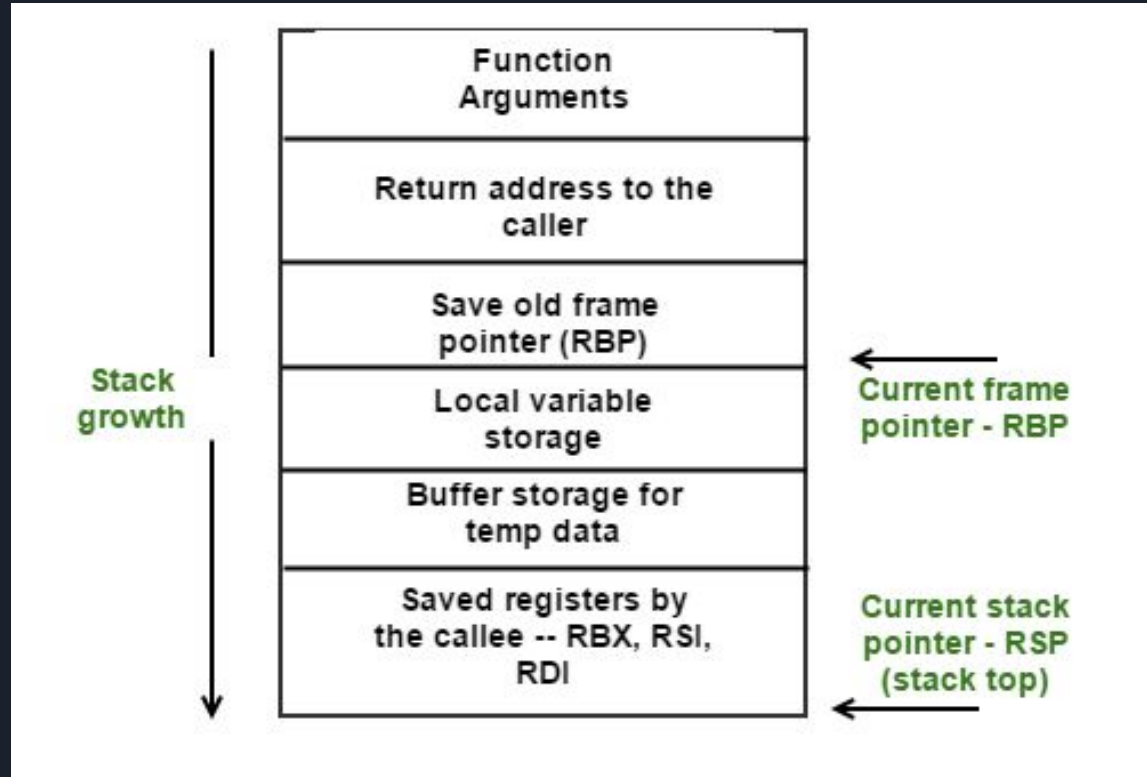


Use after free and heap corruption continue to be preferably targeted

"Other" category consists of a few common types of issues:

- XSS & zone elevation issues
- DLL planting issues
- File canonicalization & symbolic link issues

alright, so we have this stack thing...





## but what if we...


- want variables to stick around after a function call?

```
string a;  
void doSomething() {  
    a = "hello";  
}  
// A should still exist outside of doSomething()
```

- want to pick how much memory a variable uses at runtime?

```
int[] array;  
void doSomething(int size) {  
    array = new int[size]; // don't know size at runtime  
}
```

solution: keep some memory aside





# what should a heap do?

Should allow a program to:

- Ask for a chunk of memory
- Give the memory back when it's done
  - So we can give it to something else


These operations correspond to *malloc* and *free*

```
int * arr = (int *) malloc(10 * sizeof(int));  
int * arr = new int[10];
```

```
free(arr);
```

# designing a good heap is tricky

- Don't know the memory sizes program needs ahead of time
  - Need to support both small and large chunks
- Need to quickly find free memory
- Should optimally lay out memory - avoiding "fragmentation"





there is no perfect heap

Several different ways of doing it, with different tradeoffs

**Heap exploitation is taking advantage of how a heap implementation works under-the-hood.**

We will focus on the heap implementation in glibc







# basic heap implementation

- By default, all of memory is in a "wilderness" chunk
  - New memory is taken from this wilderness chunk
- The heap also maintains a list of "free" chunks
  - Chunks that have been allocated but subsequently freed
- metadata is stored next to each chunk to help with this

# how malloc works

- when we call malloc, metadata is stored nearby in headers and footers
  - int indicating size of chunk
  - flag indicating if chunk is free or in use
- extra memory padded to align the chunk, and leave room for free prev/next pointers (more on this later)



# example chunk

|                         |                                  |
|-------------------------|----------------------------------|
| an allocated<br>chunk   | size/status=inuse                |
|                         | ... user data space ...          |
|                         | size                             |
| a freed<br>chunk        | size/status=free                 |
|                         | pointer to next chunk in bin     |
|                         | pointer to previous chunk in bin |
|                         | ...unused space ...              |
|                         | size                             |
| an allocated<br>chunk   | size/status=inuse                |
|                         | user data                        |
|                         | size                             |
| other chunks            | ...                              |
| wilderness<br>chunk     | size/status=free                 |
|                         | ....                             |
|                         | size                             |
| end of available memory |                                  |


# how free chunks/bins work

- when a chunk is freed, it is placed in a linked list of freed nodes
- naive implementation: make a giant linked list of freed nodes
  - this is slow
- solution: different linked lists for chunks of various sizes

## How to keep track of free memory?

How to manage free space in the “middle” of the heap?

=> keep track of free blocks in a data structure: the “free list”. For example we could use a linked list pointing to free blocks.



A freelist!  
Good idea!

**But where do we find  
the memory to store  
this data structure?**



## Types of Bins

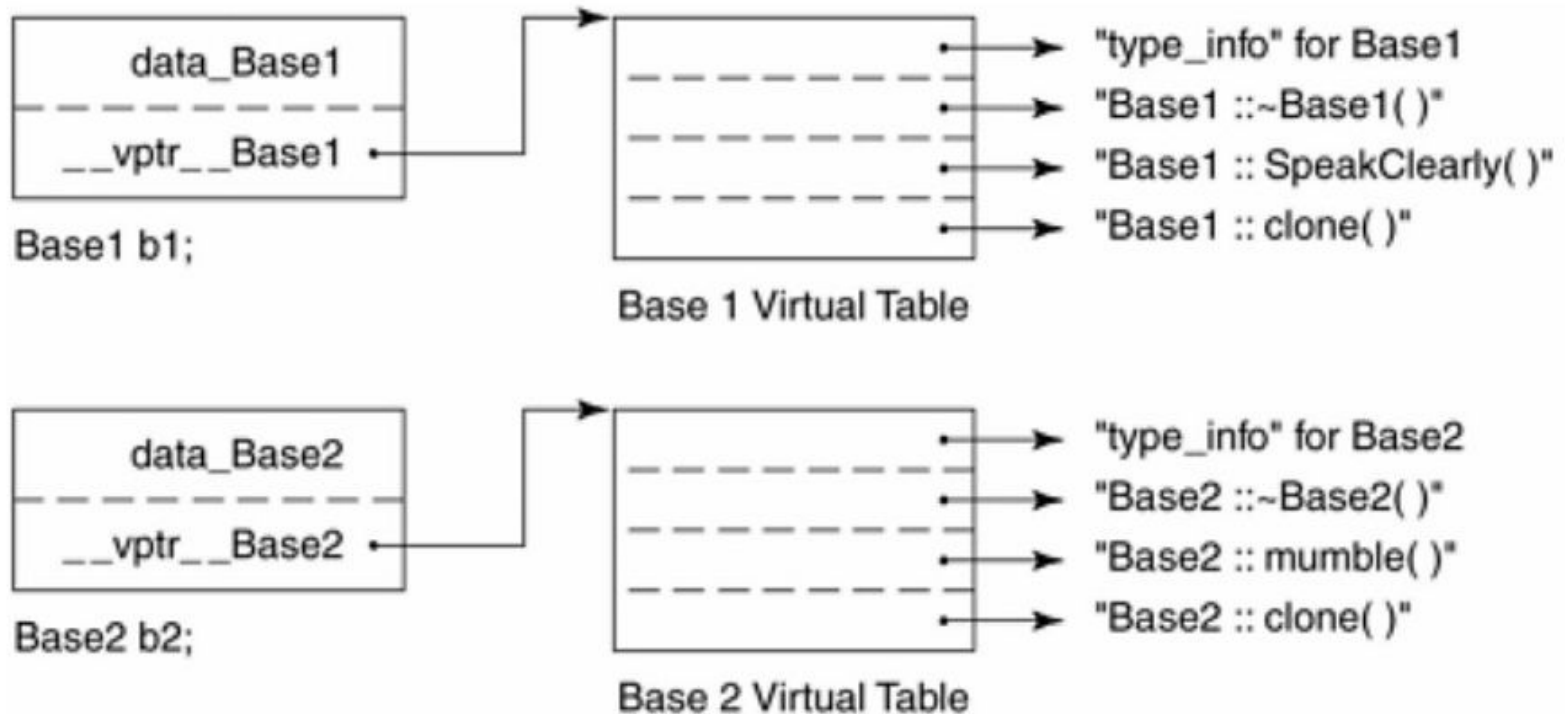
- Fast Bins: small boi, no merging, limited space
- Small Bins: Double linked list, merging
- Unsorted Bins: chunks that fit in the small bin go here
- Large Bins: big boi, sorted order, chunks within certain range, merging
- Tcache: chunks cached by current thread



# Heap Overflows

- Need to account for the structure of the heap chunks
- Common targets: function pointers, c++ vtables, etc
- Forge new chunks or corrupt adjacent chunk headers in order to fuck with the heap's free() logic

# C++ Vtables





## How to exploit?

- Overwrite the vtable pointer in our class to point to a fake vtable that we control
- Overwrite instance variables in the class

More info: <http://phrack.org/issues/56/8.html>





# Unlink Exploit

- Takes advantage of the coalescing algorithm in free() in order to write arbitrary values into any memory
- Does this by constructing fake chunks & corrupting chunks in order to mess with the free linked list pointers
- Obsolete as long as you don't use dlmalloc. Google Project Zero was able to get this exploit to work under certain circumstances: <https://googleprojectzero.blogspot.com/2014/08/the-poisoned-nul-byte-2014-edition.html>



# Double Free

- Same node added to a bin twice.
- Can be used to write values into arbitrary memory spaces

Example:

```
free(a);  Bin -> [a] -> [b] -> [a]
```

```
free(b);
```

```
free(c);
```



## Use-after-free

- Vulnerability where a freed chunk is still in use
- Similar to double free where one can call malloc to modify the contents of the chunk
- Can be used to leak information(defeat ASLR) or get arbitrary code execution
- Riskier in browsers/scripting engines(since you can reliably allocate memory on the heap) or C++ (because vtables)



## Null byte poisoning

- Triggered by an off-by-one overwrite that involves zeroing out the `PREV_IN_USE` flag in the next chunk
- Will coalesce a non free chunk, which can result in a UAF (due to overlapping chunks)



## House of Force

- requires overwrite and control over allocation size
- overwrite top chunk size chunk
- make next allocation align with libc hooks



## What to do after you get a uaf?

- overwrite the fd pointer to point to a chunk of memory to get a (nearly) arbitrary write
- leak heap and libc addresses to bypass ASLR
- unsorted bin attack(write unsorted bin linked list address anywhere)



Challenge!

shellcode.codes